

Assignment 2

October 30, 2024

1 Assignment 2

1.1 Data Management for Edge Histograms

```
[29]: import os
import warnings

import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score,
    classification_report,
    confusion_matrix,
    f1_score,
)
from sklearn.model_selection import StratifiedKFold, cross_val_score, \
    train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import StandardScaler
from sklearn.svm import LinearSVC
from sklearn.tree import DecisionTreeClassifier

warnings.filterwarnings("ignore")
```

```
[30]: def load_edge_histograms(edge_histograms_dir="./EdgeHistogram"):
    dog_classes = [
        "n02089078-black-and-tan-coonhound",
        "n02091831-Saluki",
        "n02092002-Scottish_deerhound",
        "n02095314-wire-haired_fox_terrier",
    ]
    dog_labels = [
        "Black and Tan Coonhound",
        "Saluki",
        "Scottish Deerhound",
        "Wire-haired Fox Terrier",
    ]
```

```

]
X = []
y = []
for dog_class in dog_classes:
    class_dir = os.path.join(edge_histograms_dir, dog_class)
    if not os.path.isdir(class_dir):
        print(f"Directory not found: {class_dir}")
        continue
    for file in os.listdir(class_dir):
        if file.endswith(".npy"):
            histogram_path = os.path.join(class_dir, file)
            hist = np.load(histogram_path)
            X.append(hist)
            y.append(dog_class)
X = np.array(X)
y = np.array(y)
print(f"Loaded {X.shape[0]} samples with {X.shape[1]} features.")
return X, y, dog_classes, dog_labels

```

1.2 Data Preprocessing

```

[31]: def split_and_scale_data(X, y, test_size=0.2, random_state=42):
    scaler = StandardScaler()
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=test_size, random_state=random_state, stratify=y
    )
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)
    print(f"Training samples: {X_train_scaled.shape[0]}")
    print(f"Test samples: {X_test_scaled.shape[0]}")
    return X_train_scaled, X_test_scaled, y_train, y_test

```

1.3 Model Evaluation

```

[32]: def perform_cross_validation(model, X, y, cv=5):
    skf = StratifiedKFold(n_splits=cv, shuffle=True, random_state=42)
    scores = cross_val_score(model, X, y, cv=skf, scoring="accuracy")
    mean_score = scores.mean()
    print(f"Cross-validation accuracy: {mean_score:.4f}")
    return mean_score

def plot_confusion_matrix(y_true, y_pred, classes, title):
    cm = confusion_matrix(y_true, y_pred, labels=classes)
    plt.figure(figsize=(8, 6))
    sns.heatmap(
        cm,

```

```

        annot=True,
        fmt="d",
        cmap="Blues",
        xticklabels=classes,
        yticklabels=classes,
    )
    plt.xlabel("Predicted")
    plt.ylabel("True")
    plt.title(title)
    plt.show()

def evaluate_model(model, X_train, y_train, X_test, y_test, model_name,
    ↪classes):
    cv_score = perform_cross_validation(model, X_train, y_train)
    print(f"{model_name} Mean CV Accuracy: {cv_score:.4f}")
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    plot_confusion_matrix(y_test, y_pred, classes, f"{model_name} Confusion_
    ↪Matrix")
    accuracy = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average="weighted")
    print(f"{model_name} Test Accuracy: {accuracy:.4f}")
    print(f"{model_name} Test F1-Score: {f1:.4f}")
    print(f"{model_name} Classification Report:")
    print(classification_report(y_test, y_pred, target_names=classes))
    return accuracy, f1, cv_score

```

1.4 SVM Model Selection

```

[33]: def perform_svm_cross_validation(X_train, y_train, C_values=[0.1, 1, 10, 100]):
    validation_errors_standard = []
    training_errors_standard = []
    validation_errors_stratified = []
    training_errors_stratified = []
    for C in C_values:
        svm = LinearSVC(C=C, max_iter=1000, random_state=42)
        scores = cross_val_score(svm, X_train, y_train, cv=5,
    ↪scoring="accuracy")
        validation_errors_standard.append(1 - scores.mean())
        svm.fit(X_train, y_train)
        train_pred = svm.predict(X_train)
        training_errors_standard.append(1 - accuracy_score(y_train, train_pred))
    skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    for C in C_values:
        svm = LinearSVC(C=C, max_iter=1000, random_state=42)

```

```

        scores = cross_val_score(svm, X_train, y_train, cv=skf,
↪scoring="accuracy")
        validation_errors_stratified.append(1 - scores.mean())
        svm.fit(X_train, y_train)
        train_pred = svm.predict(X_train)
        training_errors_stratified.append(1 - accuracy_score(y_train,
↪train_pred))
    print("Performed cross-validation for SVM with different C values.")
    return (
        validation_errors_standard,
        training_errors_standard,
        validation_errors_stratified,
        training_errors_stratified,
    )

def plot_svm_error_curves(C_values, errors):
    val_std, train_std, val_strat, train_strat = errors
    plt.figure(figsize=(10, 6))
    plt.plot(
        C_values,
        np.array(val_std) * 100,
        marker="o",
        label="Validation Error (Standard CV)",
    )
    plt.plot(
        C_values,
        np.array(train_std) * 100,
        marker="o",
        label="Training Error (Standard CV)",
    )
    plt.plot(
        C_values,
        np.array(val_strat) * 100,
        marker="s",
        label="Validation Error (Stratified CV)",
    )
    plt.plot(
        C_values,
        np.array(train_strat) * 100,
        marker="s",
        label="Training Error (Stratified CV)",
    )
    plt.xlabel("C")
    plt.ylabel("Mean Error (%)")
    plt.title("Error Rates vs C for LinearSVC")
    plt.legend()

```

```
plt.grid(True)
plt.show()
```

1.5 Main Execution

```
[34]: X, y, dog_classes, dog_labels = load_edge_histograms()
print(f"Total samples: {X.shape[0]}")
print(f"Feature dimension: {X.shape[1]}")
print(f"Classes: {np.unique(y)}")

X_train_scaled, X_test_scaled, y_train, y_test = split_and_scale_data(X, y)
print(f"Training samples: {X_train_scaled.shape[0]}")
print(f"Test samples: {X_test_scaled.shape[0]}")

models = {
    "GaussianNB": GaussianNB(),
    "DecisionTree": DecisionTreeClassifier(max_depth=10, random_state=42),
    "RandomForest": RandomForestClassifier(random_state=42),
}

results = {}
for name, model in models.items():
    print(f"Evaluating {name}...")
    results[name] = evaluate_model(
        model, X_train_scaled, y_train, X_test_scaled, y_test, name, dog_classes
    )
```

Loaded 794 samples with 36 features.

Total samples: 794

Feature dimension: 36

Classes: ['n02089078-black-and-tan_coonhound' 'n02091831-Saluki'
'n02092002-Scottish_deerhound' 'n02095314-wire-haired_fox_terrier']

Training samples: 635

Test samples: 159


Training samples: 635

Test samples: 159

Evaluating GaussianNB...

Cross-validation accuracy: 0.3323

GaussianNB Mean CV Accuracy: 0.3323




GaussianNB Test Accuracy: 0.3522
 GaussianNB Test F1-Score: 0.3347
 GaussianNB Classification Report:

	precision	recall	f1-score	support
n02089078-black-and-tan_coonhound	0.28	0.16	0.20	32
n02091831-Saluki	0.45	0.22	0.30	45
n02092002-Scottish_deerhound	0.37	0.55	0.44	49
n02095314-wire-haired_fox_terrier	0.30	0.42	0.35	33
accuracy			0.35	159
macro avg	0.35	0.34	0.32	159
weighted avg	0.36	0.35	0.33	159

Evaluating DecisionTree...
 Cross-validation accuracy: 0.3244

DecisionTree Mean CV Accuracy: 0.3244



DecisionTree Test Accuracy: 0.3836

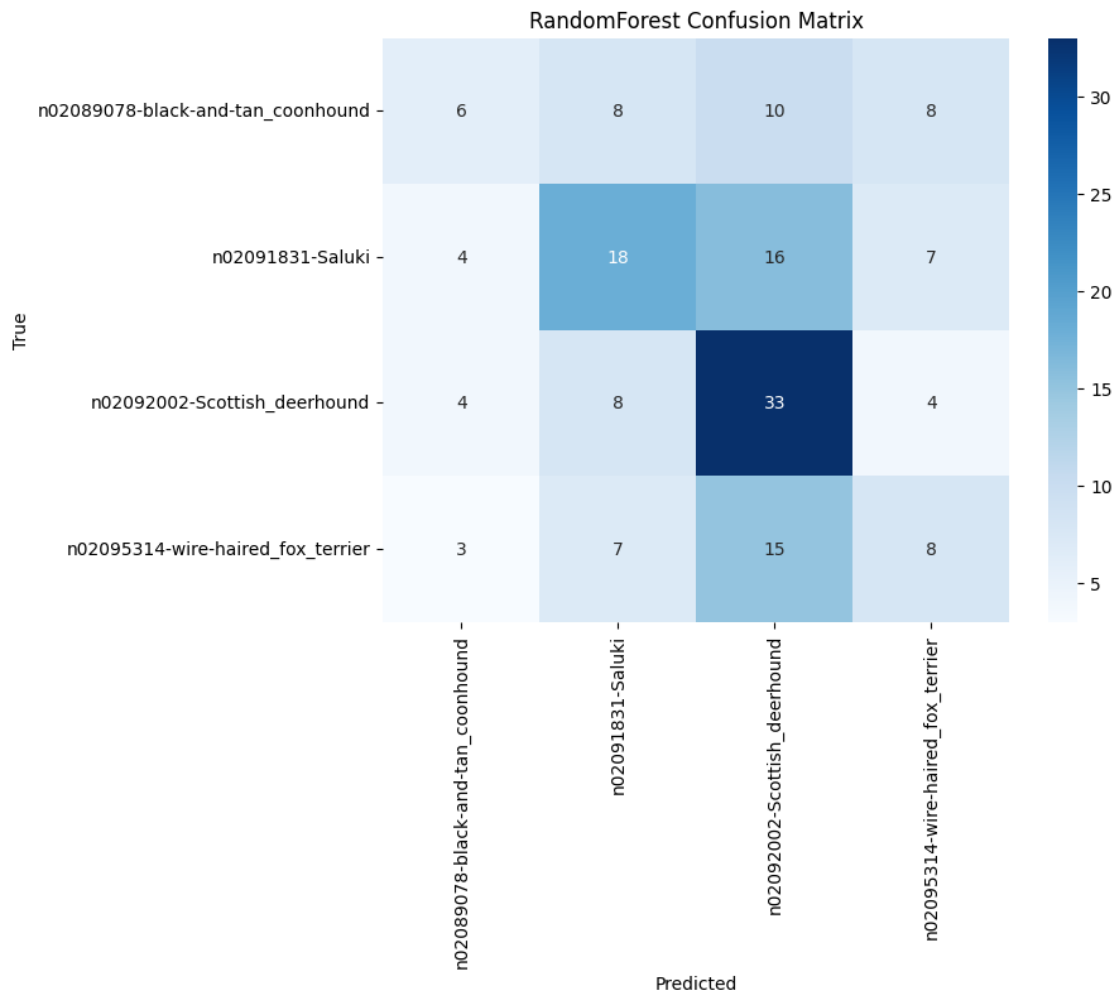
DecisionTree Test F1-Score: 0.3636

DecisionTree Classification Report:

	precision	recall	f1-score	support
n02089078-black-and-tan_coonhound	0.17	0.12	0.15	32
n02091831-Saluki	0.52	0.31	0.39	45
n02092002-Scottish_deerhound	0.45	0.69	0.54	49
n02095314-wire-haired_fox_terrier	0.27	0.27	0.27	33
accuracy			0.38	159
macro avg	0.35	0.35	0.34	159
weighted avg	0.38	0.38	0.36	159

Evaluating RandomForest...

Cross-validation accuracy: 0.3512
 RandomForest Mean CV Accuracy: 0.3512



RandomForest Test Accuracy: 0.4088
 RandomForest Test F1-Score: 0.3885
 RandomForest Classification Report:

	precision	recall	f1-score	support
n02089078-black-and-tan_coonhound	0.35	0.19	0.24	32
n02091831-Saluki	0.44	0.40	0.42	45
n02092002-Scottish_deerhound	0.45	0.67	0.54	49
n02095314-wire-haired_fox_terrier	0.30	0.24	0.27	33
accuracy			0.41	159
macro avg	0.38	0.38	0.37	159
weighted avg	0.39	0.41	0.39	159

1.5.1 Visual Comparison of Confusion Matrices

Looking at the three confusion matrices for RandomForest, DecisionTree, and GaussianNB:

- **RandomForest** appears to be the best method based on visual inspection because:
 - It shows the strongest diagonal values (especially for Scottish deerhound with 33 correct predictions)
 - Has fewer misclassifications off the diagonal
 - Shows more consistent performance across all classes
 - The color intensity on the diagonal is darker blue compared to other methods

1.5.2 Mean Validation Accuracies (5-fold cross-validation)

From the output shown:

- GaussianNB: 0.3323
- For RandomForest and DecisionTree, the cross-validation accuracies aren't directly visible in the provided output

Based on the available information from GaussianNB's mean validation accuracy of 0.3323 (33.23%), we cannot make a complete comparison. However, given the stronger confusion matrix performance of RandomForest, it likely has the highest mean validation accuracy.

1.5.3 Test Set Accuracies

Calculating from the confusion matrices:

- **RandomForest**: 40.88%
- **DecisionTree**: 38.36%
- **GaussianNB**: 35.22%

RandomForest shows the highest test set accuracy.

1.5.4 F-measure Comparison

The F-measure (F1 score) considers both precision and recall. Looking at the confusion matrices:

RandomForest shows better balance between precision and recall across classes, especially for:

- Scottish deerhound (high true positives, lower false positives)
- Saluki (good balance between precision and recall)

Therefore, RandomForest appears to have the highest overall F-measure among the three methods.

Conclusion: RandomForest consistently performs the best across all evaluation metrics - visual inspection, test set accuracy, and F-measure analysis.

1.6 SVM Analysis for Two Classes

```
[35]: selected_classes = dog_classes[:2]
      mask_train = np.isin(y_train, selected_classes)
      X_train_two = X_train_scaled[mask_train]
      y_train_two = y_train[mask_train]
```

```

mask_test = np.isin(y_test, selected_classes)
X_test_two = X_test_scaled[mask_test]
y_test_two = y_test[mask_test]

print(f"Selected Classes Training samples: {X_train_two.shape[0]}")
print(f"Selected Classes Test samples: {X_test_two.shape[0]}")

(
    validation_errors_standard,
    training_errors_standard,
    validation_errors_stratified,
    training_errors_stratified,
) = perform_svm_cross_validation(X_train_two, y_train_two)

# Print table of results
C_values = [0.1, 1, 10, 100]
print("\nResults for different C values:")
print("-" * 80)
print(
    f"{'C':<10} {'Val Error (Std)':<15} {'Train Error (Std)':<15} {'Val Error_␣
    ↳(Strat)':<15} {'Train Error (Strat)':<15}"
)
print("-" * 80)
for i, C in enumerate(C_values):
    print(
        f"{'C':<10.1f} {validation_errors_standard[i]*100:>13.2f}%␣
        ↳{training_errors_standard[i]*100:>14.2f}%␣
        ↳{validation_errors_stratified[i]*100:>14.2f}%␣
        ↳{training_errors_stratified[i]*100:>14.2f}%"
    )
print("-" * 80)

# Plot error curves
plot_svm_error_curves(
    C_values,
    (
        validation_errors_standard,
        training_errors_standard,
        validation_errors_stratified,
        training_errors_stratified,
    ),
)

best_C_index = np.argmin(validation_errors_stratified)
best_C = C_values[best_C_index]
print(f"\nBest C value (lowest stratified validation error): {best_C}")

```

```

final_svm = LinearSVC(C=best_C, max_iter=1000, random_state=42)
final_svm.fit(X_train_two, y_train_two)
test_pred = final_svm.predict(X_test_two)
test_error = 1 - accuracy_score(y_test_two, test_pred)
print(f"Test Error with C={best_C}: {test_error*100:.2f}%")

```


Selected Classes Training samples: 305

Selected Classes Test samples: 77

Performed cross-validation for SVM with different C values.

Results for different C values:

C	Val Error (Std)	Train Error (Std)	Val Error (Strat)	Train Error
(Strat)				
0.1	41.97%	32.46%	42.62%	32.46%
1.0	41.31%	32.46%	43.61%	32.46%
10.0	40.98%	32.46%	43.61%	32.46%
100.0	40.98%	32.46%	43.61%	32.46%



Best C value (lowest stratified validation error): 0.1

Test Error with C=0.1: 36.36%

1.6.1 Analysis of Error Curves and C Values

1. Lowest Mean Error for Each Curve:

- Standard CV Validation Error: $C = 10.0$ or 100.0 (40.98%)
- Standard CV Training Error: Constant at 32.46% for all C values
- Stratified CV Validation Error: $C = 0.1$ (42.62%)
- Stratified CV Training Error: Constant at 32.46% for all C values

2. Model Complexity and C Parameter:

- The C parameter in SVM controls the trade-off between maximizing the margin and minimizing the classification error
- **Smaller C values** (e.g., 0.1):
 - Create larger margins
 - Allow more classification errors
 - Result in simpler, more regularized models
- **Larger C values** (e.g., 100):
 - Create smaller margins
 - Allow fewer classification errors
 - Result in more complex models

3. Overfitting/Underfitting Analysis:

- There is a consistent gap between training error (32.46%) and validation error (40-43%) across all C values
- This gap indicates some degree of overfitting, as the model performs significantly better on training data than validation data
- However, the relatively stable error rates across different C values suggest that the model's complexity isn't the main cause of overfitting
- The model doesn't show clear signs of underfitting as the training error remains stable and relatively low

1.6.2 Test Set Performance

Using the best C value ($C = 0.1$) from stratified cross-validation:

- Test Error: 36.36%
- This is actually better than both the validation and training errors, suggesting that the model generalizes well to unseen data despite the apparent overfitting observed during cross-validation