

Assignment 3

November 26, 2024

1 Programming Assignment 3

1.1 Student ID: 916461653

```
[2]: import json
import os

import cv2
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import torch
from datasets import Dataset
from mlxtend.frequent_patterns import apriori, association_rules
from mlxtend.preprocessing import TransactionEncoder
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Conv2D, Dense, Flatten, MaxPooling2D
from tensorflow.keras.models import Sequential
from tensorflow.keras.utils import to_categorical
from torch.nn.functional import sigmoid
from transformers import (
    AutoModelForSequenceClassification,
    AutoTokenizer,
    Trainer,
    TrainingArguments,
)
```

```
/usr/local/lib/python3.10/dist-packages/tensorflow/lite/python/util.py:55:
DeprecationWarning: jax.xla_computation is deprecated. Please use the AOT APIs;
see https://jax.readthedocs.io/en/latest/aot.html. For example, replace
xla_computation(f)(*xs) with jit(f).lower(*xs).compiler_ir('hlo'). See
CHANGELOG.md for 0.4.30 for more examples.
```

```
    from jax import xla_computation as _xla_computation
/usr/local/lib/python3.10/dist-packages/tensorflow/lite/python/util.py:55:
DeprecationWarning: jax.xla_computation is deprecated. Please use the AOT APIs;
see https://jax.readthedocs.io/en/latest/aot.html. For example, replace
xla_computation(f)(*xs) with jit(f).lower(*xs).compiler_ir('hlo'). See
```

CHANGELOG.md for 0.4.30 for more examples.

```
from jax import xla_computation as _xla_computation
```

2 Question 1: Association Rule Generation from Transaction Data

In this section, we'll analyze transaction data using association rule mining. The analysis will:

1. Load and process the grocery transaction dataset
2. Calculate basic statistics about the dataset
3. Generate association rules using specified support and confidence thresholds
4. Create a heatmap visualization of rule counts for different parameter combinations

```
[ ]: print("\n" + "=" * 80)
print("PART A & B: Loading Transaction Dataset")
print("=" * 80)
print("Loading Grocery_Items_24.csv...")

transaction_data = pd.read_csv("Grocery_Items_24.csv", header=0)
grocery_transactions = transaction_data.values.tolist()
grocery_transactions = [
    [item for item in transaction if isinstance(item, str)]
    for transaction in grocery_transactions
]

print("\n" + "=" * 80)
print("PART C: Dataset Statistics")
print("=" * 80)

all_grocery_items = [
    item for transaction in grocery_transactions for item in transaction
]
unique_grocery_items = len(set(all_grocery_items))
total_transactions = len(grocery_transactions)
item_frequency = pd.Series(all_grocery_items).value_counts()
most_common_item = item_frequency.index[0]
most_common_item_count = item_frequency.iloc[0]

print(f"Number of unique items: {unique_grocery_items}")
print(f"Number of records: {total_transactions}")
print(
    f"Most popular item: {most_common_item} (appears in_
    ↪{most_common_item_count} transactions)"
)
```

```

print("\n" + "=" * 80)
print("PART D: Association Rules Generation")
print("=" * 80)
print("Generating rules with minimum support = 0.01 and minimum confidence = 0.
    ↳08...")

transaction_encoder = TransactionEncoder()
encoded_array = transaction_encoder.fit_transform(grocery_transactions)
encoded_dataframe = pd.DataFrame(encoded_array, columns=transaction_encoder.
    ↳columns_)

initial_support = 0.01
initial_confidence = 0.08
frequent_itemsets = apriori(
    encoded_dataframe, min_support=initial_support, use_colnames=True
)
initial_rules = association_rules(
    frequent_itemsets,
    frequent_itemsets,
    metric="confidence",
    min_threshold=initial_confidence,
)

print("\nAssociation Rules:")
print(initial_rules)

print("\n" + "=" * 80)
print("PART E: Heatmap Generation")
print("=" * 80)
print("Generating heatmap for different support and confidence thresholds...")

support_thresholds = [0.001, 0.005, 0.01]
confidence_thresholds = [0.05, 0.075, 0.1]
rule_count_matrix = np.zeros((len(confidence_thresholds),
    ↳len(support_thresholds)))

for confidence_idx, confidence_value in enumerate(confidence_thresholds):
    for support_idx, support_value in enumerate(support_thresholds):
        print(
            f"\nCalculating rules for support={support_value},
            ↳confidence={confidence_value}"
        )
        current_frequent_items = apriori(
            encoded_dataframe, min_support=support_value, use_colnames=True

```

```

    )
    current_rules = association_rules(
        current_frequent_items,
        current_frequent_items,
        metric="confidence",
        min_threshold=confidence_value,
    )
    rule_count_matrix[confidence_idx, support_idx] = len(current_rules)
    print(f"Number of rules found: {len(current_rules)}")

plt.figure(figsize=(10, 8))
sns.heatmap(
    rule_count_matrix,
    xticklabels=[f"{x:.3f}" for x in support_thresholds],
    yticklabels=[f"{x:.3f}" for x in confidence_thresholds],
    annot=True,
    fmt="g",
    cmap="YlOrRd",
)
plt.xlabel("Minimum Support")
plt.ylabel("Minimum Confidence")
plt.title("Number of Association Rules")
plt.tight_layout()

print("\nDisplaying heatmap...")
plt.show()

```

=====

PART A & B: Loading Transaction Dataset

=====

Loading Grocery_Items_24.csv...

=====

PART C: Dataset Statistics

=====

Number of unique items: 166

Number of records: 8000

Most popular item: whole milk (appears in 1321 transactions)

=====

PART D: Association Rules Generation

=====

Generating rules with minimum support = 0.01 and minimum confidence = 0.08...

Association Rules:

antecedents	consequents	antecedent support	\
-------------	-------------	--------------------	---

0	(rolls/buns)	(other vegetables)	0.111875
1	(other vegetables)	(rolls/buns)	0.122750
2	(whole milk)	(other vegetables)	0.156000
3	(other vegetables)	(whole milk)	0.122750
4	(whole milk)	(rolls/buns)	0.156000
5	(rolls/buns)	(whole milk)	0.111875
6	(soda)	(whole milk)	0.093625

	consequent support	support	confidence	lift	representativity \
0	0.122750	0.011250	0.100559	0.819215	1.0
1	0.111875	0.011250	0.091650	0.819215	1.0
2	0.122750	0.016125	0.103365	0.842081	1.0
3	0.156000	0.016125	0.131365	0.842081	1.0
4	0.111875	0.013500	0.086538	0.773528	1.0
5	0.156000	0.013500	0.120670	0.773528	1.0
6	0.156000	0.011375	0.121495	0.778816	1.0

	leverage	conviction	zhangs_metric	jaccard	certainty	kulczynski
0	-0.002483	0.975328	-0.199025	0.050364	-0.025296	0.096104
1	-0.002483	0.977734	-0.200997	0.050364	-0.022773	0.096104
2	-0.003024	0.978381	-0.181802	0.061399	-0.022097	0.117365
3	-0.003024	0.971639	-0.176125	0.061399	-0.029189	0.117365
4	-0.003952	0.972263	-0.257551	0.053071	-0.028528	0.103604
5	-0.003952	0.959822	-0.247927	0.053071	-0.041860	0.103604
6	-0.003231	0.960723	-0.238580	0.047744	-0.040882	0.097206

PART E: Heatmap Generation

Generating heatmap for different support and confidence thresholds...

Calculating rules for support=0.001, confidence=0.05
Number of rules found: 511

Calculating rules for support=0.005, confidence=0.05
Number of rules found: 64

Calculating rules for support=0.01, confidence=0.05
Number of rules found: 8

Calculating rules for support=0.001, confidence=0.075
Number of rules found: 294

Calculating rules for support=0.005, confidence=0.075
Number of rules found: 43


Calculating rules for support=0.01, confidence=0.075
Number of rules found: 7

Calculating rules for support=0.001, confidence=0.1
Number of rules found: 169

Calculating rules for support=0.005, confidence=0.1
Number of rules found: 27

Calculating rules for support=0.01, confidence=0.1
Number of rules found: 5

Displaying heatmap...



2.1 Dataset Characteristics (C)

- 166 unique items across 8,000 transactions
- Most frequent item: whole milk (1,321 occurrences, 16.5% of transactions)

2.2 Association Rules Analysis (D)

With minimum support=0.01 and confidence=0.08, key associations found:

1. {rolls/buns} \rightarrow {other vegetables} (conf: 10.1%, lift: 0.82)
2. {whole milk} \rightarrow {other vegetables} (conf: 10.3%, lift: 0.84)
3. {soda} \rightarrow {whole milk} (conf: 12.1%, lift: 0.78)

Lift values <1 indicate these associations occur less frequently than expected under independence.

2.3 Parameter Sensitivity Analysis (E)

The heatmap reveals:

1. Rule count decreases with increasing thresholds:
 - Max rules (511): support=0.001, confidence=0.05
 - Min rules (5): support=0.01, confidence=0.1
2. Most significant drop occurs when increasing support threshold (0.001 \rightarrow 0.005)
3. Increasing confidence has less impact on rule reduction compared to support

Optimal thresholds appear to be support=0.005, confidence=0.075 (43 rules), balancing rule quantity with significance.

3 Question 2: Image Classification using Convolutional Neural Networks

This section implements a 4-class CNN classifier with the specified architecture:

- First convolutional layer (8 3×3 filters)
- Max pooling (2×2)
- Second convolutional layer (4 3×3 filters)
- Max pooling (2×2)
- Flatten layer
- Hidden layer (8 nodes)
- Output layer (4 nodes with softmax activation)

Since my Rowan ID ends in 3, I will experiment with different numbers of nodes (4 and 16) in the hidden layer.

```
[ ]: cropped_images_dir = "./Processed"

image_data = []
image_labels = []

dog_breed_classes = [
    "n02089078-black-and-tan_coonhound",
    "n02091831-Saluki",
    "n02092002-Scottish_deerhound",
```

```

        "n02095314-wire-haired_fox_terrier",
    ]

    for breed_index, breed_class in enumerate(dog_breed_classes):
        breed_directory = os.path.join(cropped_images_dir, breed_class)
        if not os.path.isdir(breed_directory):
            print(f"Directory not found: {breed_directory}")
            continue
        for image_file in os.listdir(breed_directory):
            if image_file.endswith(".jpg"):
                image_path = os.path.join(breed_directory, image_file)
                image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
                resized_image = cv2.resize(image, (6, 6))
                image_data.append(resized_image)
                image_labels.append(breed_index)

    image_data_array = np.array(image_data)
    image_labels_array = np.array(image_labels)

    reshaped_image_data = image_data_array.reshape(-1, 6, 6, 1)
    one_hot_labels = to_categorical(image_labels_array, num_classes=4)

    X_train_initial, X_test, y_train_initial, y_test = train_test_split(
        reshaped_image_data, one_hot_labels, test_size=0.2, random_state=42
    )
    X_train, X_validation, y_train, y_validation = train_test_split(
        X_train_initial, y_train_initial, test_size=0.2, random_state=42
    )

    base_cnn_model = Sequential(
        [
            Conv2D(8, (3, 3), activation="relu", padding="same", input_shape=(6, 6, 1)),
            MaxPooling2D(pool_size=(2, 2)),
            Conv2D(4, (3, 3), activation="relu", padding="same"),
            MaxPooling2D(pool_size=(2, 2)),
            Flatten(),
            Dense(8, activation="relu"),
            Dense(4, activation="softmax"),
        ]
    )

    base_cnn_model.compile(
        optimizer="adam", loss="categorical_crossentropy", metrics=["accuracy"]
    )

```



```

)
base_model_history = base_cnn_model.fit(
    X_train,
    y_train,
    epochs=100,
    batch_size=32,
    validation_data=(X_validation, y_validation),
)

model_5x5 = Sequential(
    [
        Conv2D(8, (3, 3), activation="relu", padding="same", input_shape=(6, 6, 1)),
        MaxPooling2D(pool_size=(2, 2)),
        Conv2D(4, (5, 5), activation="relu", padding="same"),
        MaxPooling2D(pool_size=(2, 2)),
        Flatten(),
        Dense(8, activation="relu"),
        Dense(4, activation="softmax"),
    ]
)

model_5x5.compile(
    optimizer="adam", loss="categorical_crossentropy", metrics=["accuracy"]
)
model_5x5_history = model_5x5.fit(
    X_train,
    y_train,
    epochs=100,
    batch_size=32,
    validation_data=(X_validation, y_validation),
)

model_7x7 = Sequential(
    [
        Conv2D(8, (3, 3), activation="relu", padding="same", input_shape=(6, 6, 1)),
        MaxPooling2D(pool_size=(2, 2)),
        Conv2D(4, (7, 7), activation="relu", padding="same"),
        MaxPooling2D(pool_size=(2, 2)),
        Flatten(),
        Dense(8, activation="relu"),
        Dense(4, activation="softmax"),
    ]
)

```

```

model_7x7.compile(
    optimizer="adam", loss="categorical_crossentropy", metrics=["accuracy"]
)
model_7x7_history = model_7x7.fit(
    X_train,
    y_train,
    epochs=100,
    batch_size=32,
    validation_data=(X_validation, y_validation),
)

plt.figure(figsize=(12, 4))

plt.subplot(1, 3, 1)
plt.plot(base_model_history.history["accuracy"], label="Training Accuracy")
plt.plot(base_model_history.history["val_accuracy"], label="Validation Accuracy")
plt.title("Base Model (3x3 filter)")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()

plt.subplot(1, 3, 2)
plt.plot(model_5x5_history.history["accuracy"], label="Training Accuracy")
plt.plot(model_5x5_history.history["val_accuracy"], label="Validation Accuracy")
plt.title("Model with 5x5 filter")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()

plt.subplot(1, 3, 3)
plt.plot(model_7x7_history.history["accuracy"], label="Training Accuracy")
plt.plot(model_7x7_history.history["val_accuracy"], label="Validation Accuracy")
plt.title("Model with 7x7 filter")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()

plt.tight_layout()
plt.show()

print("\nFinal Accuracies:")
print(
    f"Base Model (3x3) - Training: {base_model_history.history['accuracy'][-1]:.4f}, Validation: {base_model_history.history['val_accuracy'][-1]:.4f}"
)

```

```

)
print(
    f"5x5 Filter Model - Training: {model_5x5_history.history['accuracy'][-1]:.4f}, Validation: {model_5x5_history.history['val_accuracy'][-1]:.4f}"
)
print(
    f"7x7 Filter Model - Training: {model_7x7_history.history['accuracy'][-1]:.4f}, Validation: {model_7x7_history.history['val_accuracy'][-1]:.4f}"
)

```

Epoch 1/100

/Users/rudra/Documents/Code/Freelancing/Python/Data Mining/.venv/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

16/16 0s 5ms/step - accuracy: 0.2853 - loss: 15.3956 - val_accuracy: 0.2520 - val_loss: 5.1526
Epoch 2/100

16/16 0s 1ms/step - accuracy: 0.2256 - loss: 4.2694 - val_accuracy: 0.1575 - val_loss: 1.9203
Epoch 3/100

16/16 0s 1ms/step - accuracy: 0.2404 - loss: 1.9045 - val_accuracy: 0.2520 - val_loss: 1.5187
Epoch 4/100

16/16 0s 1ms/step - accuracy: 0.2640 - loss: 1.5503 - val_accuracy: 0.2677 - val_loss: 1.4922
Epoch 5/100

16/16 0s 1ms/step - accuracy: 0.2533 - loss: 1.4825 - val_accuracy: 0.2677 - val_loss: 1.4777
Epoch 6/100

16/16 0s 1ms/step - accuracy: 0.2596 - loss: 1.4559 - val_accuracy: 0.2598 - val_loss: 1.4682
Epoch 7/100

16/16 0s 1ms/step - accuracy: 0.2629 - loss: 1.4389 - val_accuracy: 0.2598 - val_loss: 1.4609
Epoch 8/100

16/16 0s 1ms/step - accuracy: 0.2610 - loss: 1.4261 - val_accuracy: 0.2598 - val_loss: 1.4546
Epoch 9/100

16/16 0s 1ms/step - accuracy: 0.2628 - loss: 1.4212 - val_accuracy: 0.2520 - val_loss: 1.4485
Epoch 10/100

16/16 0s 1ms/step -

accuracy: 0.2628 - loss: 1.4106 - val_accuracy: 0.2520 - val_loss: 1.4421
 Epoch 11/100
 16/16 0s 1ms/step -
 accuracy: 0.2649 - loss: 1.4006 - val_accuracy: 0.2520 - val_loss: 1.4371
 Epoch 12/100
 16/16 0s 4ms/step -
 accuracy: 0.2690 - loss: 1.3938 - val_accuracy: 0.2520 - val_loss: 1.4333
 Epoch 13/100
 16/16 0s 1ms/step -
 accuracy: 0.2638 - loss: 1.3904 - val_accuracy: 0.2520 - val_loss: 1.4304
 Epoch 14/100
 16/16 0s 1ms/step -
 accuracy: 0.2638 - loss: 1.3895 - val_accuracy: 0.2520 - val_loss: 1.4279
 Epoch 15/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3885 - val_accuracy: 0.2598 - val_loss: 1.4254
 Epoch 16/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3875 - val_accuracy: 0.2598 - val_loss: 1.4235
 Epoch 17/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3865 - val_accuracy: 0.2598 - val_loss: 1.4217
 Epoch 18/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3855 - val_accuracy: 0.2598 - val_loss: 1.4200
 Epoch 19/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3846 - val_accuracy: 0.2598 - val_loss: 1.4184
 Epoch 20/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3838 - val_accuracy: 0.2598 - val_loss: 1.4170
 Epoch 21/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3830 - val_accuracy: 0.2598 - val_loss: 1.4157
 Epoch 22/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3822 - val_accuracy: 0.2598 - val_loss: 1.4144
 Epoch 23/100
 16/16 0s 1ms/step -
 accuracy: 0.2640 - loss: 1.3815 - val_accuracy: 0.2598 - val_loss: 1.4132
 Epoch 24/100
 16/16 0s 4ms/step -
 accuracy: 0.2640 - loss: 1.3809 - val_accuracy: 0.2598 - val_loss: 1.4120
 Epoch 25/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3803 - val_accuracy: 0.2598 - val_loss: 1.4109
 Epoch 26/100
 16/16 0s 1ms/step -

accuracy: 0.2658 - loss: 1.3800 - val_accuracy: 0.2598 - val_loss: 1.4101
 Epoch 27/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3799 - val_accuracy: 0.2598 - val_loss: 1.4094
 Epoch 28/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3797 - val_accuracy: 0.2598 - val_loss: 1.4088
 Epoch 29/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3795 - val_accuracy: 0.2598 - val_loss: 1.4082
 Epoch 30/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3794 - val_accuracy: 0.2598 - val_loss: 1.4074
 Epoch 31/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3793 - val_accuracy: 0.2598 - val_loss: 1.4068
 Epoch 32/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3792 - val_accuracy: 0.2598 - val_loss: 1.4064
 Epoch 33/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3791 - val_accuracy: 0.2598 - val_loss: 1.4059
 Epoch 34/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3790 - val_accuracy: 0.2598 - val_loss: 1.4055
 Epoch 35/100
 16/16 0s 1ms/step -
 accuracy: 0.2651 - loss: 1.3790 - val_accuracy: 0.2598 - val_loss: 1.4051
 Epoch 36/100
 16/16 0s 4ms/step -
 accuracy: 0.2658 - loss: 1.3789 - val_accuracy: 0.2598 - val_loss: 1.4049
 Epoch 37/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3788 - val_accuracy: 0.2598 - val_loss: 1.4045
 Epoch 38/100
 16/16 0s 1ms/step -
 accuracy: 0.2670 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4042
 Epoch 39/100
 16/16 0s 1ms/step -
 accuracy: 0.2670 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4036
 Epoch 40/100
 16/16 0s 1ms/step -
 accuracy: 0.2670 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4032
 Epoch 41/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4029
 Epoch 42/100
 16/16 0s 1ms/step -

accuracy: 0.2670 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4026
 Epoch 43/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4024
 Epoch 44/100
 16/16 0s 1ms/step -
 accuracy: 0.2651 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4020
 Epoch 45/100
 16/16 0s 1ms/step -
 accuracy: 0.2658 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4017
 Epoch 46/100
 16/16 0s 1ms/step -
 accuracy: 0.2670 - loss: 1.3788 - val_accuracy: 0.2677 - val_loss: 1.4016
 Epoch 47/100
 16/16 0s 4ms/step -
 accuracy: 0.2750 - loss: 1.3788 - val_accuracy: 0.3622 - val_loss: 1.4014
 Epoch 48/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3788 - val_accuracy: 0.3622 - val_loss: 1.4013
 Epoch 49/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3789 - val_accuracy: 0.3622 - val_loss: 1.4010
 Epoch 50/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3789 - val_accuracy: 0.3622 - val_loss: 1.4008
 Epoch 51/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3789 - val_accuracy: 0.3622 - val_loss: 1.4009
 Epoch 52/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3789 - val_accuracy: 0.3622 - val_loss: 1.4008
 Epoch 53/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3789 - val_accuracy: 0.3622 - val_loss: 1.4008
 Epoch 54/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3789 - val_accuracy: 0.3622 - val_loss: 1.4008
 Epoch 55/100
 16/16 0s 3ms/step -
 accuracy: 0.2950 - loss: 1.3791 - val_accuracy: 0.3622 - val_loss: 1.4002
 Epoch 56/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3790 - val_accuracy: 0.3622 - val_loss: 1.4001
 Epoch 57/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3790 - val_accuracy: 0.3622 - val_loss: 1.4002
 Epoch 58/100
 16/16 0s 1ms/step -

accuracy: 0.2950 - loss: 1.3790 - val_accuracy: 0.3622 - val_loss: 1.4008
 Epoch 59/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3791 - val_accuracy: 0.3622 - val_loss: 1.4002
 Epoch 60/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3791 - val_accuracy: 0.3622 - val_loss: 1.3998
 Epoch 61/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3791 - val_accuracy: 0.3622 - val_loss: 1.3999
 Epoch 62/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3792 - val_accuracy: 0.3622 - val_loss: 1.4000
 Epoch 63/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3792 - val_accuracy: 0.3622 - val_loss: 1.4000
 Epoch 64/100
 16/16 0s 2ms/step -
 accuracy: 0.2950 - loss: 1.3792 - val_accuracy: 0.3622 - val_loss: 1.4005
 Epoch 65/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3792 - val_accuracy: 0.3622 - val_loss: 1.4005
 Epoch 66/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.3998
 Epoch 67/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.4001
 Epoch 68/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.3996
 Epoch 69/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.3996
 Epoch 70/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.4000
 Epoch 71/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.4011
 Epoch 72/100
 16/16 0s 2ms/step -
 accuracy: 0.2950 - loss: 1.3795 - val_accuracy: 0.3622 - val_loss: 1.3997
 Epoch 73/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3794 - val_accuracy: 0.3622 - val_loss: 1.3997
 Epoch 74/100
 16/16 0s 2ms/step -

accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.4005
 Epoch 75/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.4008
 Epoch 76/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3793 - val_accuracy: 0.3622 - val_loss: 1.4009
 Epoch 77/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3795 - val_accuracy: 0.3622 - val_loss: 1.4011
 Epoch 78/100
 16/16 0s 2ms/step -
 accuracy: 0.2956 - loss: 1.3791 - val_accuracy: 0.3622 - val_loss: 1.4011
 Epoch 79/100
 16/16 0s 2ms/step -
 accuracy: 0.2950 - loss: 1.3791 - val_accuracy: 0.3622 - val_loss: 1.4022
 Epoch 80/100
 16/16 0s 2ms/step -
 accuracy: 0.2950 - loss: 1.3792 - val_accuracy: 0.3622 - val_loss: 1.4010
 Epoch 81/100
 16/16 0s 2ms/step -
 accuracy: 0.2950 - loss: 1.3791 - val_accuracy: 0.3622 - val_loss: 1.4018
 Epoch 82/100
 16/16 0s 2ms/step -
 accuracy: 0.2950 - loss: 1.3789 - val_accuracy: 0.3622 - val_loss: 1.4027
 Epoch 83/100
 16/16 0s 4ms/step -
 accuracy: 0.2950 - loss: 1.3790 - val_accuracy: 0.3622 - val_loss: 1.4017
 Epoch 84/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3789 - val_accuracy: 0.3622 - val_loss: 1.4015
 Epoch 85/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3788 - val_accuracy: 0.3622 - val_loss: 1.4025
 Epoch 86/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3787 - val_accuracy: 0.3622 - val_loss: 1.4032
 Epoch 87/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3788 - val_accuracy: 0.3622 - val_loss: 1.4028
 Epoch 88/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3786 - val_accuracy: 0.3622 - val_loss: 1.4041
 Epoch 89/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3786 - val_accuracy: 0.3622 - val_loss: 1.4039
 Epoch 90/100
 16/16 0s 2ms/step -

accuracy: 0.2950 - loss: 1.3786 - val_accuracy: 0.3622 - val_loss: 1.4031
 Epoch 91/100
 16/16 0s 1ms/step -
 accuracy: 0.2952 - loss: 1.3785 - val_accuracy: 0.3622 - val_loss: 1.4040
 Epoch 92/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3784 - val_accuracy: 0.3622 - val_loss: 1.4043
 Epoch 93/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3779 - val_accuracy: 0.3622 - val_loss: 1.4035
 Epoch 94/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3777 - val_accuracy: 0.3622 - val_loss: 1.4041
 Epoch 95/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3771 - val_accuracy: 0.3622 - val_loss: 1.4050
 Epoch 96/100
 16/16 0s 4ms/step -
 accuracy: 0.2950 - loss: 1.3767 - val_accuracy: 0.3622 - val_loss: 1.4067
 Epoch 97/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3764 - val_accuracy: 0.3622 - val_loss: 1.4050
 Epoch 98/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3766 - val_accuracy: 0.3622 - val_loss: 1.4063
 Epoch 99/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3760 - val_accuracy: 0.3622 - val_loss: 1.4075
 Epoch 100/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3761 - val_accuracy: 0.3622 - val_loss: 1.4073
 Epoch 1/100
 16/16 0s 5ms/step -
 accuracy: 0.2445 - loss: 11.0896 - val_accuracy: 0.3701 - val_loss: 3.2967
 Epoch 2/100
 16/16 0s 2ms/step -
 accuracy: 0.3389 - loss: 2.2577 - val_accuracy: 0.2756 - val_loss: 1.3954
 Epoch 3/100
 16/16 0s 3ms/step -
 accuracy: 0.2636 - loss: 1.4034 - val_accuracy: 0.2835 - val_loss: 1.3868
 Epoch 4/100
 16/16 0s 1ms/step -
 accuracy: 0.2611 - loss: 1.3881 - val_accuracy: 0.2756 - val_loss: 1.3865
 Epoch 5/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3869 - val_accuracy: 0.2756 - val_loss: 1.3846
 Epoch 6/100
 16/16 0s 1ms/step -

accuracy: 0.2656 - loss: 1.3859 - val_accuracy: 0.2756 - val_loss: 1.3829
 Epoch 7/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3851 - val_accuracy: 0.2756 - val_loss: 1.3815
 Epoch 8/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3844 - val_accuracy: 0.2756 - val_loss: 1.3800
 Epoch 9/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3839 - val_accuracy: 0.2756 - val_loss: 1.3789
 Epoch 10/100
 16/16 0s 2ms/step -
 accuracy: 0.2656 - loss: 1.3834 - val_accuracy: 0.2756 - val_loss: 1.3782
 Epoch 11/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3830 - val_accuracy: 0.2756 - val_loss: 1.3772
 Epoch 12/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3826 - val_accuracy: 0.2756 - val_loss: 1.3762
 Epoch 13/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3822 - val_accuracy: 0.2756 - val_loss: 1.3756
 Epoch 14/100
 16/16 0s 2ms/step -
 accuracy: 0.2656 - loss: 1.3819 - val_accuracy: 0.2756 - val_loss: 1.3751
 Epoch 15/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3816 - val_accuracy: 0.2756 - val_loss: 1.3745
 Epoch 16/100
 16/16 0s 4ms/step -
 accuracy: 0.2656 - loss: 1.3813 - val_accuracy: 0.2756 - val_loss: 1.3738
 Epoch 17/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3811 - val_accuracy: 0.2756 - val_loss: 1.3730
 Epoch 18/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3809 - val_accuracy: 0.2756 - val_loss: 1.3722
 Epoch 19/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3806 - val_accuracy: 0.2756 - val_loss: 1.3713
 Epoch 20/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3804 - val_accuracy: 0.2756 - val_loss: 1.3708
 Epoch 21/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3803 - val_accuracy: 0.2835 - val_loss: 1.3696
 Epoch 22/100
 16/16 0s 5ms/step -

accuracy: 0.2656 - loss: 1.3801 - val_accuracy: 0.2835 - val_loss: 1.3684
 Epoch 23/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3798 - val_accuracy: 0.2835 - val_loss: 1.3676
 Epoch 24/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3796 - val_accuracy: 0.2835 - val_loss: 1.3670
 Epoch 25/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3795 - val_accuracy: 0.2835 - val_loss: 1.3664
 Epoch 26/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3796 - val_accuracy: 0.2835 - val_loss: 1.3662
 Epoch 27/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3793 - val_accuracy: 0.2835 - val_loss: 1.3656
 Epoch 28/100
 16/16 0s 4ms/step -
 accuracy: 0.2637 - loss: 1.3792 - val_accuracy: 0.2835 - val_loss: 1.3654
 Epoch 29/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3792 - val_accuracy: 0.2835 - val_loss: 1.3650
 Epoch 30/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3790 - val_accuracy: 0.2835 - val_loss: 1.3641
 Epoch 31/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3789 - val_accuracy: 0.2835 - val_loss: 1.3633
 Epoch 32/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3788 - val_accuracy: 0.2835 - val_loss: 1.3628
 Epoch 33/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3787 - val_accuracy: 0.2835 - val_loss: 1.3629
 Epoch 34/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3788 - val_accuracy: 0.2835 - val_loss: 1.3629
 Epoch 35/100
 16/16 0s 2ms/step -
 accuracy: 0.2637 - loss: 1.3786 - val_accuracy: 0.2835 - val_loss: 1.3618
 Epoch 36/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3786 - val_accuracy: 0.2835 - val_loss: 1.3613
 Epoch 37/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3785 - val_accuracy: 0.2835 - val_loss: 1.3612
 Epoch 38/100
 16/16 0s 1ms/step -

accuracy: 0.2637 - loss: 1.3784 - val_accuracy: 0.2835 - val_loss: 1.3609
 Epoch 39/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3784 - val_accuracy: 0.2835 - val_loss: 1.3605
 Epoch 40/100
 16/16 0s 1ms/step -
 accuracy: 0.2637 - loss: 1.3782 - val_accuracy: 0.2835 - val_loss: 1.3609
 Epoch 41/100
 16/16 0s 4ms/step -
 accuracy: 0.2637 - loss: 1.3783 - val_accuracy: 0.2835 - val_loss: 1.3603
 Epoch 42/100
 16/16 0s 2ms/step -
 accuracy: 0.2656 - loss: 1.3781 - val_accuracy: 0.2835 - val_loss: 1.3600
 Epoch 43/100
 16/16 0s 1ms/step -
 accuracy: 0.2662 - loss: 1.3780 - val_accuracy: 0.2835 - val_loss: 1.3599
 Epoch 44/100
 16/16 0s 1ms/step -
 accuracy: 0.2656 - loss: 1.3780 - val_accuracy: 0.2835 - val_loss: 1.3592
 Epoch 45/100
 16/16 0s 1ms/step -
 accuracy: 0.2591 - loss: 1.3779 - val_accuracy: 0.2835 - val_loss: 1.3588
 Epoch 46/100
 16/16 0s 1ms/step -
 accuracy: 0.2598 - loss: 1.3778 - val_accuracy: 0.3701 - val_loss: 1.3592
 Epoch 47/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3778 - val_accuracy: 0.3701 - val_loss: 1.3584
 Epoch 48/100
 16/16 0s 5ms/step -
 accuracy: 0.2997 - loss: 1.3777 - val_accuracy: 0.3701 - val_loss: 1.3581
 Epoch 49/100
 16/16 0s 1ms/step -
 accuracy: 0.3003 - loss: 1.3775 - val_accuracy: 0.3701 - val_loss: 1.3586
 Epoch 50/100
 16/16 0s 1ms/step -
 accuracy: 0.2950 - loss: 1.3774 - val_accuracy: 0.3701 - val_loss: 1.3577
 Epoch 51/100
 16/16 0s 1ms/step -
 accuracy: 0.3003 - loss: 1.3774 - val_accuracy: 0.3701 - val_loss: 1.3584
 Epoch 52/100
 16/16 0s 1ms/step -
 accuracy: 0.2904 - loss: 1.3773 - val_accuracy: 0.3701 - val_loss: 1.3571
 Epoch 53/100
 16/16 0s 2ms/step -
 accuracy: 0.3003 - loss: 1.3772 - val_accuracy: 0.3701 - val_loss: 1.3573
 Epoch 54/100
 16/16 0s 1ms/step -

accuracy: 0.2950 - loss: 1.3771 - val_accuracy: 0.3701 - val_loss: 1.3574
 Epoch 55/100
 16/16 0s 2ms/step -
 accuracy: 0.3003 - loss: 1.3770 - val_accuracy: 0.3701 - val_loss: 1.3569
 Epoch 56/100
 16/16 0s 5ms/step -
 accuracy: 0.3003 - loss: 1.3769 - val_accuracy: 0.3701 - val_loss: 1.3580
 Epoch 57/100
 16/16 0s 2ms/step -
 accuracy: 0.2930 - loss: 1.3767 - val_accuracy: 0.3701 - val_loss: 1.3588
 Epoch 58/100
 16/16 0s 2ms/step -
 accuracy: 0.2909 - loss: 1.3769 - val_accuracy: 0.3701 - val_loss: 1.3563
 Epoch 59/100
 16/16 0s 2ms/step -
 accuracy: 0.2952 - loss: 1.3765 - val_accuracy: 0.3701 - val_loss: 1.3572
 Epoch 60/100
 16/16 0s 2ms/step -
 accuracy: 0.2909 - loss: 1.3765 - val_accuracy: 0.3701 - val_loss: 1.3551
 Epoch 61/100
 16/16 0s 1ms/step -
 accuracy: 0.2978 - loss: 1.3766 - val_accuracy: 0.3622 - val_loss: 1.3552
 Epoch 62/100
 16/16 0s 1ms/step -
 accuracy: 0.2983 - loss: 1.3762 - val_accuracy: 0.3701 - val_loss: 1.3561
 Epoch 63/100
 16/16 0s 1ms/step -
 accuracy: 0.2901 - loss: 1.3762 - val_accuracy: 0.3701 - val_loss: 1.3538
 Epoch 64/100
 16/16 0s 2ms/step -
 accuracy: 0.2978 - loss: 1.3764 - val_accuracy: 0.3622 - val_loss: 1.3537
 Epoch 65/100
 16/16 0s 2ms/step -
 accuracy: 0.2952 - loss: 1.3759 - val_accuracy: 0.3701 - val_loss: 1.3567
 Epoch 66/100
 16/16 0s 5ms/step -
 accuracy: 0.2834 - loss: 1.3758 - val_accuracy: 0.3701 - val_loss: 1.3537
 Epoch 67/100
 16/16 0s 2ms/step -
 accuracy: 0.2962 - loss: 1.3755 - val_accuracy: 0.3701 - val_loss: 1.3550
 Epoch 68/100
 16/16 0s 2ms/step -
 accuracy: 0.2862 - loss: 1.3753 - val_accuracy: 0.3701 - val_loss: 1.3558
 Epoch 69/100
 16/16 0s 1ms/step -
 accuracy: 0.2841 - loss: 1.3747 - val_accuracy: 0.3701 - val_loss: 1.3556
 Epoch 70/100
 16/16 0s 1ms/step -

accuracy: 0.2848 - loss: 1.3738 - val_accuracy: 0.3701 - val_loss: 1.3541
 Epoch 71/100
 16/16 0s 2ms/step -
 accuracy: 0.2877 - loss: 1.3735 - val_accuracy: 0.3858 - val_loss: 1.3605
 Epoch 72/100
 16/16 0s 2ms/step -
 accuracy: 0.2939 - loss: 1.3748 - val_accuracy: 0.3780 - val_loss: 1.3557
 Epoch 73/100
 16/16 0s 2ms/step -
 accuracy: 0.2862 - loss: 1.3712 - val_accuracy: 0.3858 - val_loss: 1.3575
 Epoch 74/100
 16/16 0s 1ms/step -
 accuracy: 0.2953 - loss: 1.3697 - val_accuracy: 0.3622 - val_loss: 1.3584
 Epoch 75/100
 16/16 0s 1ms/step -
 accuracy: 0.3055 - loss: 1.3666 - val_accuracy: 0.3622 - val_loss: 1.3597
 Epoch 76/100
 16/16 0s 1ms/step -
 accuracy: 0.3087 - loss: 1.3641 - val_accuracy: 0.3937 - val_loss: 1.3535
 Epoch 77/100
 16/16 0s 1ms/step -
 accuracy: 0.3033 - loss: 1.3559 - val_accuracy: 0.4016 - val_loss: 1.3457
 Epoch 78/100
 16/16 0s 5ms/step -
 accuracy: 0.3139 - loss: 1.3479 - val_accuracy: 0.4173 - val_loss: 1.3366
 Epoch 79/100
 16/16 0s 1ms/step -
 accuracy: 0.3173 - loss: 1.3387 - val_accuracy: 0.4331 - val_loss: 1.3243
 Epoch 80/100
 16/16 0s 2ms/step -
 accuracy: 0.3293 - loss: 1.3258 - val_accuracy: 0.4252 - val_loss: 1.3168
 Epoch 81/100
 16/16 0s 2ms/step -
 accuracy: 0.3572 - loss: 1.3148 - val_accuracy: 0.4331 - val_loss: 1.3112
 Epoch 82/100
 16/16 0s 1ms/step -
 accuracy: 0.3599 - loss: 1.3085 - val_accuracy: 0.4331 - val_loss: 1.3085
 Epoch 83/100
 16/16 0s 1ms/step -
 accuracy: 0.3499 - loss: 1.3010 - val_accuracy: 0.4409 - val_loss: 1.3050
 Epoch 84/100
 16/16 0s 1ms/step -
 accuracy: 0.3515 - loss: 1.3018 - val_accuracy: 0.4173 - val_loss: 1.2958
 Epoch 85/100
 16/16 0s 1ms/step -
 accuracy: 0.3506 - loss: 1.2908 - val_accuracy: 0.4173 - val_loss: 1.2865
 Epoch 86/100
 16/16 0s 1ms/step -

accuracy: 0.3572 - loss: 1.2829 - val_accuracy: 0.4331 - val_loss: 1.2779
 Epoch 87/100
 16/16 0s 2ms/step -
 accuracy: 0.3659 - loss: 1.2635 - val_accuracy: 0.4488 - val_loss: 1.2728
 Epoch 88/100
 16/16 0s 1ms/step -
 accuracy: 0.3664 - loss: 1.2581 - val_accuracy: 0.4252 - val_loss: 1.2612
 Epoch 89/100
 16/16 0s 2ms/step -
 accuracy: 0.3615 - loss: 1.2544 - val_accuracy: 0.4331 - val_loss: 1.2655
 Epoch 90/100
 16/16 0s 3ms/step -
 accuracy: 0.3538 - loss: 1.2431 - val_accuracy: 0.4409 - val_loss: 1.2626
 Epoch 91/100
 16/16 0s 1ms/step -
 accuracy: 0.3538 - loss: 1.2356 - val_accuracy: 0.4409 - val_loss: 1.2581
 Epoch 92/100
 16/16 0s 1ms/step -
 accuracy: 0.3598 - loss: 1.2278 - val_accuracy: 0.4488 - val_loss: 1.2581
 Epoch 93/100
 16/16 0s 2ms/step -
 accuracy: 0.3463 - loss: 1.2271 - val_accuracy: 0.4409 - val_loss: 1.2559
 Epoch 94/100
 16/16 0s 2ms/step -
 accuracy: 0.3461 - loss: 1.2215 - val_accuracy: 0.4488 - val_loss: 1.2569
 Epoch 95/100
 16/16 0s 1ms/step -
 accuracy: 0.3661 - loss: 1.2155 - val_accuracy: 0.4567 - val_loss: 1.2599
 Epoch 96/100
 16/16 0s 1ms/step -
 accuracy: 0.3485 - loss: 1.2212 - val_accuracy: 0.4567 - val_loss: 1.2525
 Epoch 97/100
 16/16 0s 2ms/step -
 accuracy: 0.3701 - loss: 1.2078 - val_accuracy: 0.4567 - val_loss: 1.2569
 Epoch 98/100
 16/16 0s 2ms/step -
 accuracy: 0.3592 - loss: 1.2016 - val_accuracy: 0.4567 - val_loss: 1.2585
 Epoch 99/100
 16/16 0s 2ms/step -
 accuracy: 0.3578 - loss: 1.2040 - val_accuracy: 0.4567 - val_loss: 1.2546
 Epoch 100/100
 16/16 0s 1ms/step -
 accuracy: 0.3573 - loss: 1.1959 - val_accuracy: 0.4488 - val_loss: 1.2561
 Epoch 1/100
 16/16 0s 6ms/step -
 accuracy: 0.2358 - loss: 4.4041 - val_accuracy: 0.2283 - val_loss: 3.1856
 Epoch 2/100
 16/16 0s 2ms/step -

accuracy: 0.2714 - loss: 2.7035 - val_accuracy: 0.2677 - val_loss: 2.6461
 Epoch 3/100
 16/16 0s 2ms/step -
 accuracy: 0.2963 - loss: 2.2935 - val_accuracy: 0.2598 - val_loss: 2.2401
 Epoch 4/100
 16/16 0s 2ms/step -
 accuracy: 0.2667 - loss: 2.0322 - val_accuracy: 0.3071 - val_loss: 2.0417
 Epoch 5/100
 16/16 0s 2ms/step -
 accuracy: 0.2769 - loss: 1.8951 - val_accuracy: 0.2756 - val_loss: 1.8776
 Epoch 6/100
 16/16 0s 2ms/step -
 accuracy: 0.2488 - loss: 1.7771 - val_accuracy: 0.2835 - val_loss: 1.7231
 Epoch 7/100
 16/16 0s 5ms/step -
 accuracy: 0.2568 - loss: 1.5950 - val_accuracy: 0.2992 - val_loss: 1.6612
 Epoch 8/100
 16/16 0s 2ms/step -
 accuracy: 0.2831 - loss: 1.4712 - val_accuracy: 0.3071 - val_loss: 1.6225
 Epoch 9/100
 16/16 0s 2ms/step -
 accuracy: 0.3021 - loss: 1.4120 - val_accuracy: 0.3307 - val_loss: 1.5567
 Epoch 10/100
 16/16 0s 2ms/step -
 accuracy: 0.3122 - loss: 1.3497 - val_accuracy: 0.3386 - val_loss: 1.5118
 Epoch 11/100
 16/16 0s 2ms/step -
 accuracy: 0.3571 - loss: 1.2914 - val_accuracy: 0.3307 - val_loss: 1.4641
 Epoch 12/100
 16/16 0s 2ms/step -
 accuracy: 0.3663 - loss: 1.2607 - val_accuracy: 0.2913 - val_loss: 1.4477
 Epoch 13/100
 16/16 0s 2ms/step -
 accuracy: 0.3894 - loss: 1.2443 - val_accuracy: 0.2992 - val_loss: 1.4299
 Epoch 14/100
 16/16 0s 2ms/step -
 accuracy: 0.3841 - loss: 1.2120 - val_accuracy: 0.3150 - val_loss: 1.4191
 Epoch 15/100
 16/16 0s 2ms/step -
 accuracy: 0.4104 - loss: 1.1830 - val_accuracy: 0.2913 - val_loss: 1.4119
 Epoch 16/100
 16/16 0s 2ms/step -
 accuracy: 0.4126 - loss: 1.1611 - val_accuracy: 0.2835 - val_loss: 1.4055
 Epoch 17/100
 16/16 0s 2ms/step -
 accuracy: 0.4317 - loss: 1.1419 - val_accuracy: 0.2913 - val_loss: 1.3967
 Epoch 18/100
 16/16 0s 2ms/step -

accuracy: 0.4327 - loss: 1.1285 - val_accuracy: 0.2677 - val_loss: 1.4014
 Epoch 19/100
 16/16 0s 3ms/step -
 accuracy: 0.4522 - loss: 1.1134 - val_accuracy: 0.2677 - val_loss: 1.3965
 Epoch 20/100
 16/16 0s 2ms/step -
 accuracy: 0.4650 - loss: 1.0965 - val_accuracy: 0.2756 - val_loss: 1.3943
 Epoch 21/100
 16/16 0s 2ms/step -
 accuracy: 0.4838 - loss: 1.0865 - val_accuracy: 0.2835 - val_loss: 1.3959
 Epoch 22/100
 16/16 0s 2ms/step -
 accuracy: 0.4727 - loss: 1.0764 - val_accuracy: 0.2913 - val_loss: 1.3876
 Epoch 23/100
 16/16 0s 2ms/step -
 accuracy: 0.4902 - loss: 1.0630 - val_accuracy: 0.3228 - val_loss: 1.3739
 Epoch 24/100
 16/16 0s 2ms/step -
 accuracy: 0.5046 - loss: 1.0504 - val_accuracy: 0.3465 - val_loss: 1.3792
 Epoch 25/100
 16/16 0s 2ms/step -
 accuracy: 0.5145 - loss: 1.0342 - val_accuracy: 0.3307 - val_loss: 1.3965
 Epoch 26/100
 16/16 0s 2ms/step -
 accuracy: 0.5289 - loss: 1.0258 - val_accuracy: 0.3307 - val_loss: 1.3994
 Epoch 27/100
 16/16 0s 2ms/step -
 accuracy: 0.5170 - loss: 1.0197 - val_accuracy: 0.3150 - val_loss: 1.4060
 Epoch 28/100
 16/16 0s 2ms/step -
 accuracy: 0.5277 - loss: 1.0104 - val_accuracy: 0.2992 - val_loss: 1.4132
 Epoch 29/100
 16/16 0s 2ms/step -
 accuracy: 0.5375 - loss: 1.0022 - val_accuracy: 0.3228 - val_loss: 1.4182
 Epoch 30/100
 16/16 0s 3ms/step -
 accuracy: 0.5453 - loss: 0.9956 - val_accuracy: 0.3228 - val_loss: 1.4235
 Epoch 31/100
 16/16 0s 2ms/step -
 accuracy: 0.5475 - loss: 0.9859 - val_accuracy: 0.3228 - val_loss: 1.4271
 Epoch 32/100
 16/16 0s 2ms/step -
 accuracy: 0.5491 - loss: 0.9771 - val_accuracy: 0.3228 - val_loss: 1.4368
 Epoch 33/100
 16/16 0s 2ms/step -
 accuracy: 0.5623 - loss: 0.9684 - val_accuracy: 0.3307 - val_loss: 1.4419
 Epoch 34/100
 16/16 0s 2ms/step -


accuracy: 0.5666 - loss: 0.9606 - val_accuracy: 0.3307 - val_loss: 1.4434
 Epoch 35/100
 16/16 0s 3ms/step -
 accuracy: 0.5658 - loss: 0.9543 - val_accuracy: 0.3307 - val_loss: 1.4550
 Epoch 36/100
 16/16 0s 2ms/step -
 accuracy: 0.5681 - loss: 0.9503 - val_accuracy: 0.3150 - val_loss: 1.4490
 Epoch 37/100
 16/16 0s 5ms/step -
 accuracy: 0.5650 - loss: 0.9431 - val_accuracy: 0.3307 - val_loss: 1.4200
 Epoch 38/100
 16/16 0s 2ms/step -
 accuracy: 0.5791 - loss: 0.9335 - val_accuracy: 0.3543 - val_loss: 1.4109
 Epoch 39/100
 16/16 0s 2ms/step -
 accuracy: 0.5954 - loss: 0.9274 - val_accuracy: 0.3465 - val_loss: 1.4185
 Epoch 40/100
 16/16 0s 2ms/step -
 accuracy: 0.5897 - loss: 0.9211 - val_accuracy: 0.3465 - val_loss: 1.4316
 Epoch 41/100
 16/16 0s 2ms/step -
 accuracy: 0.5999 - loss: 0.9147 - val_accuracy: 0.3386 - val_loss: 1.4409
 Epoch 42/100
 16/16 0s 2ms/step -
 accuracy: 0.6120 - loss: 0.9140 - val_accuracy: 0.3465 - val_loss: 1.4497
 Epoch 43/100
 16/16 0s 4ms/step -
 accuracy: 0.6014 - loss: 0.9065 - val_accuracy: 0.3780 - val_loss: 1.4673
 Epoch 44/100
 16/16 0s 5ms/step -
 accuracy: 0.6112 - loss: 0.9023 - val_accuracy: 0.3858 - val_loss: 1.4814
 Epoch 45/100
 16/16 0s 2ms/step -
 accuracy: 0.6036 - loss: 0.9001 - val_accuracy: 0.3780 - val_loss: 1.4905
 Epoch 46/100
 16/16 0s 1ms/step -
 accuracy: 0.6171 - loss: 0.8962 - val_accuracy: 0.3622 - val_loss: 1.4953
 Epoch 47/100
 16/16 0s 2ms/step -
 accuracy: 0.6069 - loss: 0.8964 - val_accuracy: 0.3622 - val_loss: 1.5025
 Epoch 48/100
 16/16 0s 2ms/step -
 accuracy: 0.6112 - loss: 0.8885 - val_accuracy: 0.3622 - val_loss: 1.5269
 Epoch 49/100
 16/16 0s 2ms/step -
 accuracy: 0.5990 - loss: 0.8879 - val_accuracy: 0.3622 - val_loss: 1.5139
 Epoch 50/100
 16/16 0s 2ms/step -

accuracy: 0.6183 - loss: 0.8778 - val_accuracy: 0.3622 - val_loss: 1.5086
 Epoch 51/100
 16/16 0s 2ms/step -
 accuracy: 0.6222 - loss: 0.8785 - val_accuracy: 0.3701 - val_loss: 1.4817
 Epoch 52/100
 16/16 0s 1ms/step -
 accuracy: 0.6266 - loss: 0.8665 - val_accuracy: 0.3701 - val_loss: 1.4823
 Epoch 53/100
 16/16 0s 1ms/step -
 accuracy: 0.6115 - loss: 0.8645 - val_accuracy: 0.3858 - val_loss: 1.4495
 Epoch 54/100
 16/16 0s 1ms/step -
 accuracy: 0.6299 - loss: 0.8699 - val_accuracy: 0.4094 - val_loss: 1.4457
 Epoch 55/100
 16/16 0s 2ms/step -
 accuracy: 0.6219 - loss: 0.8687 - val_accuracy: 0.4331 - val_loss: 1.4418
 Epoch 56/100
 16/16 0s 1ms/step -
 accuracy: 0.6133 - loss: 0.8770 - val_accuracy: 0.4567 - val_loss: 1.4481
 Epoch 57/100
 16/16 0s 1ms/step -
 accuracy: 0.6036 - loss: 0.8774 - val_accuracy: 0.4331 - val_loss: 1.4611
 Epoch 58/100
 16/16 0s 4ms/step -
 accuracy: 0.6082 - loss: 0.8821 - val_accuracy: 0.4488 - val_loss: 1.4773
 Epoch 59/100
 16/16 0s 1ms/step -
 accuracy: 0.6031 - loss: 0.8929 - val_accuracy: 0.4331 - val_loss: 1.4254
 Epoch 60/100
 16/16 0s 2ms/step -
 accuracy: 0.6163 - loss: 0.8708 - val_accuracy: 0.4488 - val_loss: 1.4233
 Epoch 61/100
 16/16 0s 2ms/step -
 accuracy: 0.6202 - loss: 0.8550 - val_accuracy: 0.4331 - val_loss: 1.4274
 Epoch 62/100
 16/16 0s 1ms/step -
 accuracy: 0.6286 - loss: 0.8465 - val_accuracy: 0.4252 - val_loss: 1.4497
 Epoch 63/100
 16/16 0s 1ms/step -
 accuracy: 0.6367 - loss: 0.8360 - val_accuracy: 0.4252 - val_loss: 1.4750
 Epoch 64/100
 16/16 0s 1ms/step -
 accuracy: 0.6243 - loss: 0.8327 - val_accuracy: 0.4016 - val_loss: 1.4833
 Epoch 65/100
 16/16 0s 2ms/step -
 accuracy: 0.6333 - loss: 0.8324 - val_accuracy: 0.4094 - val_loss: 1.4973
 Epoch 66/100
 16/16 0s 4ms/step -

accuracy: 0.6357 - loss: 0.8303 - val_accuracy: 0.4094 - val_loss: 1.4764
Epoch 67/100
16/16 0s 1ms/step -
accuracy: 0.6378 - loss: 0.8269 - val_accuracy: 0.4331 - val_loss: 1.4806
Epoch 68/100
16/16 0s 1ms/step -
accuracy: 0.6463 - loss: 0.8084 - val_accuracy: 0.4252 - val_loss: 1.4877
Epoch 69/100
16/16 0s 1ms/step -
accuracy: 0.6561 - loss: 0.8006 - val_accuracy: 0.4173 - val_loss: 1.4840
Epoch 70/100
16/16 0s 1ms/step -
accuracy: 0.6714 - loss: 0.7842 - val_accuracy: 0.4094 - val_loss: 1.5025
Epoch 71/100
16/16 0s 2ms/step -
accuracy: 0.6812 - loss: 0.7734 - val_accuracy: 0.4331 - val_loss: 1.5098
Epoch 72/100
16/16 0s 1ms/step -
accuracy: 0.6731 - loss: 0.7670 - val_accuracy: 0.3937 - val_loss: 1.5154
Epoch 73/100
16/16 0s 4ms/step -
accuracy: 0.6629 - loss: 0.7630 - val_accuracy: 0.3937 - val_loss: 1.5241
Epoch 74/100
16/16 0s 1ms/step -
accuracy: 0.6796 - loss: 0.7606 - val_accuracy: 0.3701 - val_loss: 1.5465
Epoch 75/100
16/16 0s 1ms/step -
accuracy: 0.6896 - loss: 0.7541 - val_accuracy: 0.3780 - val_loss: 1.5406
Epoch 76/100
16/16 0s 2ms/step -
accuracy: 0.6849 - loss: 0.7527 - val_accuracy: 0.3780 - val_loss: 1.5656
Epoch 77/100
16/16 0s 2ms/step -
accuracy: 0.6938 - loss: 0.7433 - val_accuracy: 0.3701 - val_loss: 1.5541
Epoch 78/100
16/16 0s 2ms/step -
accuracy: 0.6885 - loss: 0.7443 - val_accuracy: 0.3780 - val_loss: 1.5856
Epoch 79/100
16/16 0s 1ms/step -
accuracy: 0.6924 - loss: 0.7384 - val_accuracy: 0.3858 - val_loss: 1.5692
Epoch 80/100
16/16 0s 4ms/step -
accuracy: 0.6930 - loss: 0.7395 - val_accuracy: 0.3780 - val_loss: 1.5666
Epoch 81/100
16/16 0s 2ms/step -
accuracy: 0.7026 - loss: 0.7330 - val_accuracy: 0.3937 - val_loss: 1.5447
Epoch 82/100
16/16 0s 1ms/step -

accuracy: 0.7007 - loss: 0.7334 - val_accuracy: 0.3701 - val_loss: 1.5764
Epoch 83/100
16/16 0s 2ms/step -
accuracy: 0.6827 - loss: 0.7299 - val_accuracy: 0.3858 - val_loss: 1.5687
Epoch 84/100
16/16 0s 1ms/step -
accuracy: 0.6924 - loss: 0.7274 - val_accuracy: 0.4016 - val_loss: 1.5885
Epoch 85/100
16/16 0s 2ms/step -
accuracy: 0.6897 - loss: 0.7308 - val_accuracy: 0.3937 - val_loss: 1.5856
Epoch 86/100
16/16 0s 2ms/step -
accuracy: 0.7031 - loss: 0.7244 - val_accuracy: 0.3858 - val_loss: 1.6176
Epoch 87/100
16/16 0s 3ms/step -
accuracy: 0.6999 - loss: 0.7249 - val_accuracy: 0.3858 - val_loss: 1.6101
Epoch 88/100
16/16 0s 1ms/step -
accuracy: 0.6871 - loss: 0.7211 - val_accuracy: 0.4173 - val_loss: 1.5464
Epoch 89/100
16/16 0s 1ms/step -
accuracy: 0.7069 - loss: 0.7292 - val_accuracy: 0.4331 - val_loss: 1.5485
Epoch 90/100
16/16 0s 1ms/step -
accuracy: 0.7175 - loss: 0.7260 - val_accuracy: 0.4173 - val_loss: 1.5467
Epoch 91/100
16/16 0s 1ms/step -
accuracy: 0.7086 - loss: 0.7162 - val_accuracy: 0.4331 - val_loss: 1.5431
Epoch 92/100
16/16 0s 1ms/step -
accuracy: 0.6994 - loss: 0.7247 - val_accuracy: 0.4724 - val_loss: 1.5275
Epoch 93/100
16/16 0s 1ms/step -
accuracy: 0.6941 - loss: 0.7279 - val_accuracy: 0.4646 - val_loss: 1.5378
Epoch 94/100
16/16 0s 4ms/step -
accuracy: 0.7008 - loss: 0.7228 - val_accuracy: 0.4803 - val_loss: 1.5408
Epoch 95/100
16/16 0s 1ms/step -
accuracy: 0.6826 - loss: 0.7314 - val_accuracy: 0.4724 - val_loss: 1.5701
Epoch 96/100
16/16 0s 2ms/step -
accuracy: 0.6902 - loss: 0.7309 - val_accuracy: 0.4803 - val_loss: 1.5679
Epoch 97/100
16/16 0s 2ms/step -
accuracy: 0.6954 - loss: 0.7364 - val_accuracy: 0.4646 - val_loss: 1.5731
Epoch 98/100
16/16 0s 2ms/step -

accuracy: 0.6820 - loss: 0.7410 - val_accuracy: 0.4567 - val_loss: 1.5536
Epoch 99/100
16/16 0s 1ms/step -
accuracy: 0.6812 - loss: 0.7325 - val_accuracy: 0.4724 - val_loss: 1.5493
Epoch 100/100
16/16 0s 1ms/step -
accuracy: 0.6762 - loss: 0.7279 - val_accuracy: 0.4488 - val_loss: 1.5328



Final Accuracies:

Base Model (3x3) - Training: 0.3012, Validation: 0.3622
5x5 Filter Model - Training: 0.3760, Validation: 0.4488
7x7 Filter Model - Training: 0.6713, Validation: 0.4488

3.1 Model Architecture & Performance

3.1.1 Base Model (3×3 filter)

- Training: 0.3012, Validation: 0.3622
- Shows consistent but slow improvement over epochs
- Higher validation than training suggests underfitting
- Stable learning curve with good generalization
- Limited by small receptive field for large features

3.1.2 5×5 Filter Model

- Training: 0.3760, Validation: 0.4488
- Performance improvement from base model
- Shows unstable validation accuracy
- Training curve shows limited learning capacity
- Filter size possibly too large for 6×6 input images

3.1.3 7×7 Filter Model

- Training: 0.6713, Validation: 0.4488
- Best overall performance

- Strongest learning progression
- Shows significant overfitting (train > val)
- Captures global image features effectively

3.2 Model Fitness Analysis

1. Base Model (3×3):

- Underfits the data
- Receptive field too small for global patterns
- Good stability but insufficient learning capacity

2. 5×5 Model:

- Moderate underfitting
- Filter size mismatched with input dimensions
- Better feature extraction than base model

3. 7×7 Model:

- Shows significant overfitting
- Highest training accuracy but validation matches 5x5
- Most effective at capturing training patterns

3.3 Key Finding

Larger filter size (7×7) demonstrates superior training performance but overfits, while the 5×5 filter achieves the same validation accuracy with better generalization. This suggests that intermediate filter sizes may provide better balance between feature capture and model generalization for small-resolution images.

4 Question 3: Text Classification using BERT

This section implements a multi-label text classification system by fine-tuning BERT. The task involves:

- Processing tweet data with 11 emotion classes
- Fine-tuning BERT base model for multi-label classification
- Evaluating performance using both strict and flexible matching criteria
- Visualizing the learning process through training curves

```
[3]: computation_device = torch.device("cuda" if torch.cuda.is_available() else
    ↪ "cpu")
    print(f"Using device: {computation_device}")

    emotion_labels = [
        "anger",
        "anticipation",
        "disgust",
```

```

    "fear",
    "joy",
    "love",
    "optimism",
    "pessimism",
    "sadness",
    "surprise",
    "trust",
]

emotion_index_to_label = {idx: label for idx, label in
    ↪enumerate(emotion_labels)}
emotion_label_to_index = {label: idx for idx, label in
    ↪enumerate(emotion_labels)}

print("Loading datasets...")
training_data = [json.loads(line) for line in open("./train.json", "r")]
validation_data = [json.loads(line) for line in open("./validation.json", "r")]
testing_data = [json.loads(line) for line in open("./test.json", "r")]

training_dataframe = pd.DataFrame(training_data)
validation_dataframe = pd.DataFrame(validation_data)
testing_dataframe = pd.DataFrame(testing_data)

print("Initializing tokenizer...")
bert_tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

print("Converting to HuggingFace datasets...")
training_dataset = Dataset.from_pandas(training_dataframe)
validation_dataset = Dataset.from_pandas(validation_dataframe)
testing_dataset = Dataset.from_pandas(testing_dataframe)

print("Preprocessing datasets...")
tokenized_training = bert_tokenizer(
    training_dataset["Tweet"], padding="max_length", truncation=True,
    ↪max_length=128
)
training_label_matrix = np.zeros((len(training_dataset["Tweet"]),
    ↪len(emotion_labels)))
for idx, label in enumerate(emotion_labels):
    training_label_matrix[:, idx] = training_dataset[label]
tokenized_training["labels"] = training_label_matrix.tolist()
training_dataset = Dataset.from_dict(tokenized_training)

tokenized_validation = bert_tokenizer(
    validation_dataset["Tweet"], padding="max_length", truncation=True,
    ↪max_length=128

```



```

)
validation_label_matrix = np.zeros(
    (len(validation_dataset["Tweet"]), len(emotion_labels))
)
for idx, label in enumerate(emotion_labels):
    validation_label_matrix[:, idx] = validation_dataset[label]
tokenized_validation["labels"] = validation_label_matrix.tolist()
validation_dataset = Dataset.from_dict(tokenized_validation)

tokenized_testing = bert_tokenizer(
    testing_dataset["Tweet"], padding="max_length", truncation=True,
    ↪max_length=128
)
testing_label_matrix = np.zeros((len(testing_dataset["Tweet"]),
    ↪len(emotion_labels)))
for idx, label in enumerate(emotion_labels):
    testing_label_matrix[:, idx] = testing_dataset[label]
tokenized_testing["labels"] = testing_label_matrix.tolist()
testing_dataset = Dataset.from_dict(tokenized_testing)

training_dataset.set_format("torch")
validation_dataset.set_format("torch")
testing_dataset.set_format("torch")

print("Initializing model...")
emotion_classifier = AutoModelForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    problem_type="multi_label_classification",
    num_labels=len(emotion_labels),
    id2label=emotion_index_to_label,
    label2id=emotion_label_to_index,
)

print("Setting up training arguments...")
training_configuration = TrainingArguments(
    output_dir="./bert_output",
    learning_rate=2e-5,
    per_device_train_batch_size=8,
    per_device_eval_batch_size=8,
    num_train_epochs=5,
    weight_decay=0.01,
    evaluation_strategy="epoch",
    save_strategy="epoch",
    load_best_model_at_end=True,
    metric_for_best_model="accuracy",
    logging_dir="./logs",
    logging_strategy="steps",

```

```

        logging_steps=10,
        remove_unused_columns=False,
        report_to="none",
        save_total_limit=2,
    )

    print("Initializing trainer...")
    model_trainer = Trainer(
        model=emotion_classifier,
        args=training_configuration,
        train_dataset=training_dataset,
        eval_dataset=validation_dataset,
        compute_metrics=lambda eval_pred: {
            "accuracy": accuracy_score(
                (sigmoid(torch.tensor(eval_pred[0])).numpy() > 0.5).astype(np.
float32),
                eval_pred[1],
            )
        },
    )

    print("Starting training...")
    training_results = model_trainer.train()

    training_history = model_trainer.state.log_history
    training_metrics = [
        (log["epoch"], log["loss"])
        for log in training_history
        if "loss" in log and "eval_loss" not in log
    ]
    validation_metrics = [
        (log["epoch"], log["eval_loss"]) for log in training_history if "eval_loss"
in log
    ]
    training_metrics.sort(key=lambda x: x[0])
    validation_metrics.sort(key=lambda x: x[0])
    training_epochs, training_losses = zip(*training_metrics)
    validation_epochs, validation_losses = zip(*validation_metrics)

    plt.figure(figsize=(10, 6))
    plt.plot(training_epochs, training_losses, "b-", label="Training Loss")
    plt.plot(validation_epochs, validation_losses, "r-", label="Validation Loss")
    plt.title("Training and Validation Loss Curves")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.grid(True)

```

```

plt.xticks(range(0, int(max(training_epochs)) + 1))
plt.close()

print("\nEvaluating with strict accuracy...")
strict_test_results = model_trainer.evaluate(testing_dataset)
print("\nTest Results (Strict Accuracy - all labels must match):")
print(f"Accuracy: {strict_test_results['eval_accuracy']:.4f}")

model_trainer.compute_metrics = lambda eval_pred: {
    "accuracy": (
        (sigmoid(torch.tensor(eval_pred[0])).numpy() > 0.5).astype(np.float32)
        == eval_pred[1]
    )
    .any(axis=1)
    .mean()
}
any_match_test_results = model_trainer.evaluate(testing_dataset)
print("\nTest Results (Any-Match Accuracy - at least one label must match):")
print(f"Accuracy: {any_match_test_results['eval_accuracy']:.4f}")

```

```

/usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283:
DeprecationWarning: `should_run_async` will not call `transform_cell`
automatically in the future. Please pass the result to `transformed_cell`
argument and any exception that happen during the transform in
`preprocessing_exc_tuple` in IPython 7.17 and above.
    and should_run_async(code)

```

Using device: cuda

Loading datasets...

Initializing tokenizer...

```

/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_auth.py:94:

```

UserWarning:

The secret `HF_TOKEN` does not exist in your Colab secrets.

To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as secret in your Google Colab and restart your session.

You will be able to reuse this secret in all of your notebooks.

Please note that authentication is recommended but still optional to access public models or datasets.

```
warnings.warn(
```

```
tokenizer_config.json: 0%|          | 0.00/48.0 [00:00<?, ?B/s]
```

```
config.json: 0%|          | 0.00/570 [00:00<?, ?B/s]
```

```
vocab.txt: 0%|          | 0.00/232k [00:00<?, ?B/s]
```

```
tokenizer.json: 0%|          | 0.00/466k [00:00<?, ?B/s]
```

Converting to HuggingFace datasets...

Preprocessing datasets...

Initializing model...

model.safetensors: 0%| | 0.00/440M [00:00<?, ?B/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:

['classifier.bias', 'classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

/usr/local/lib/python3.10/dist-packages/transformers/training_args.py:1568:

FutureWarning: `evaluation_strategy` is deprecated and will be removed in version 4.46 of Transformers. Use `eval_strategy` instead

warnings.warn(

Setting up training arguments...

Initializing trainer...

Starting training...

<IPython.core.display.HTML object>

Evaluating with strict accuracy...

<IPython.core.display.HTML object>

Test Results (Strict Accuracy - all labels must match):

Accuracy: 0.2580

Test Results (Any-Match Accuracy - at least one label must match):

Accuracy: 1.0000